

# Natural Language Processing

## Week 1

# Agenda

- Instructor & Student Introductions
- Course Introduction
- Regular Expressions
- Text Normalization
- Minimum Edit Distance

# About the Instructor

## Professional Experience

5 years experience in data analytics

Applied ML in 2 domains: finance, law

## Education

MS in Applied Data Science, University of Chicago

BS in Physics, University of Chicago

CFA Charterholder

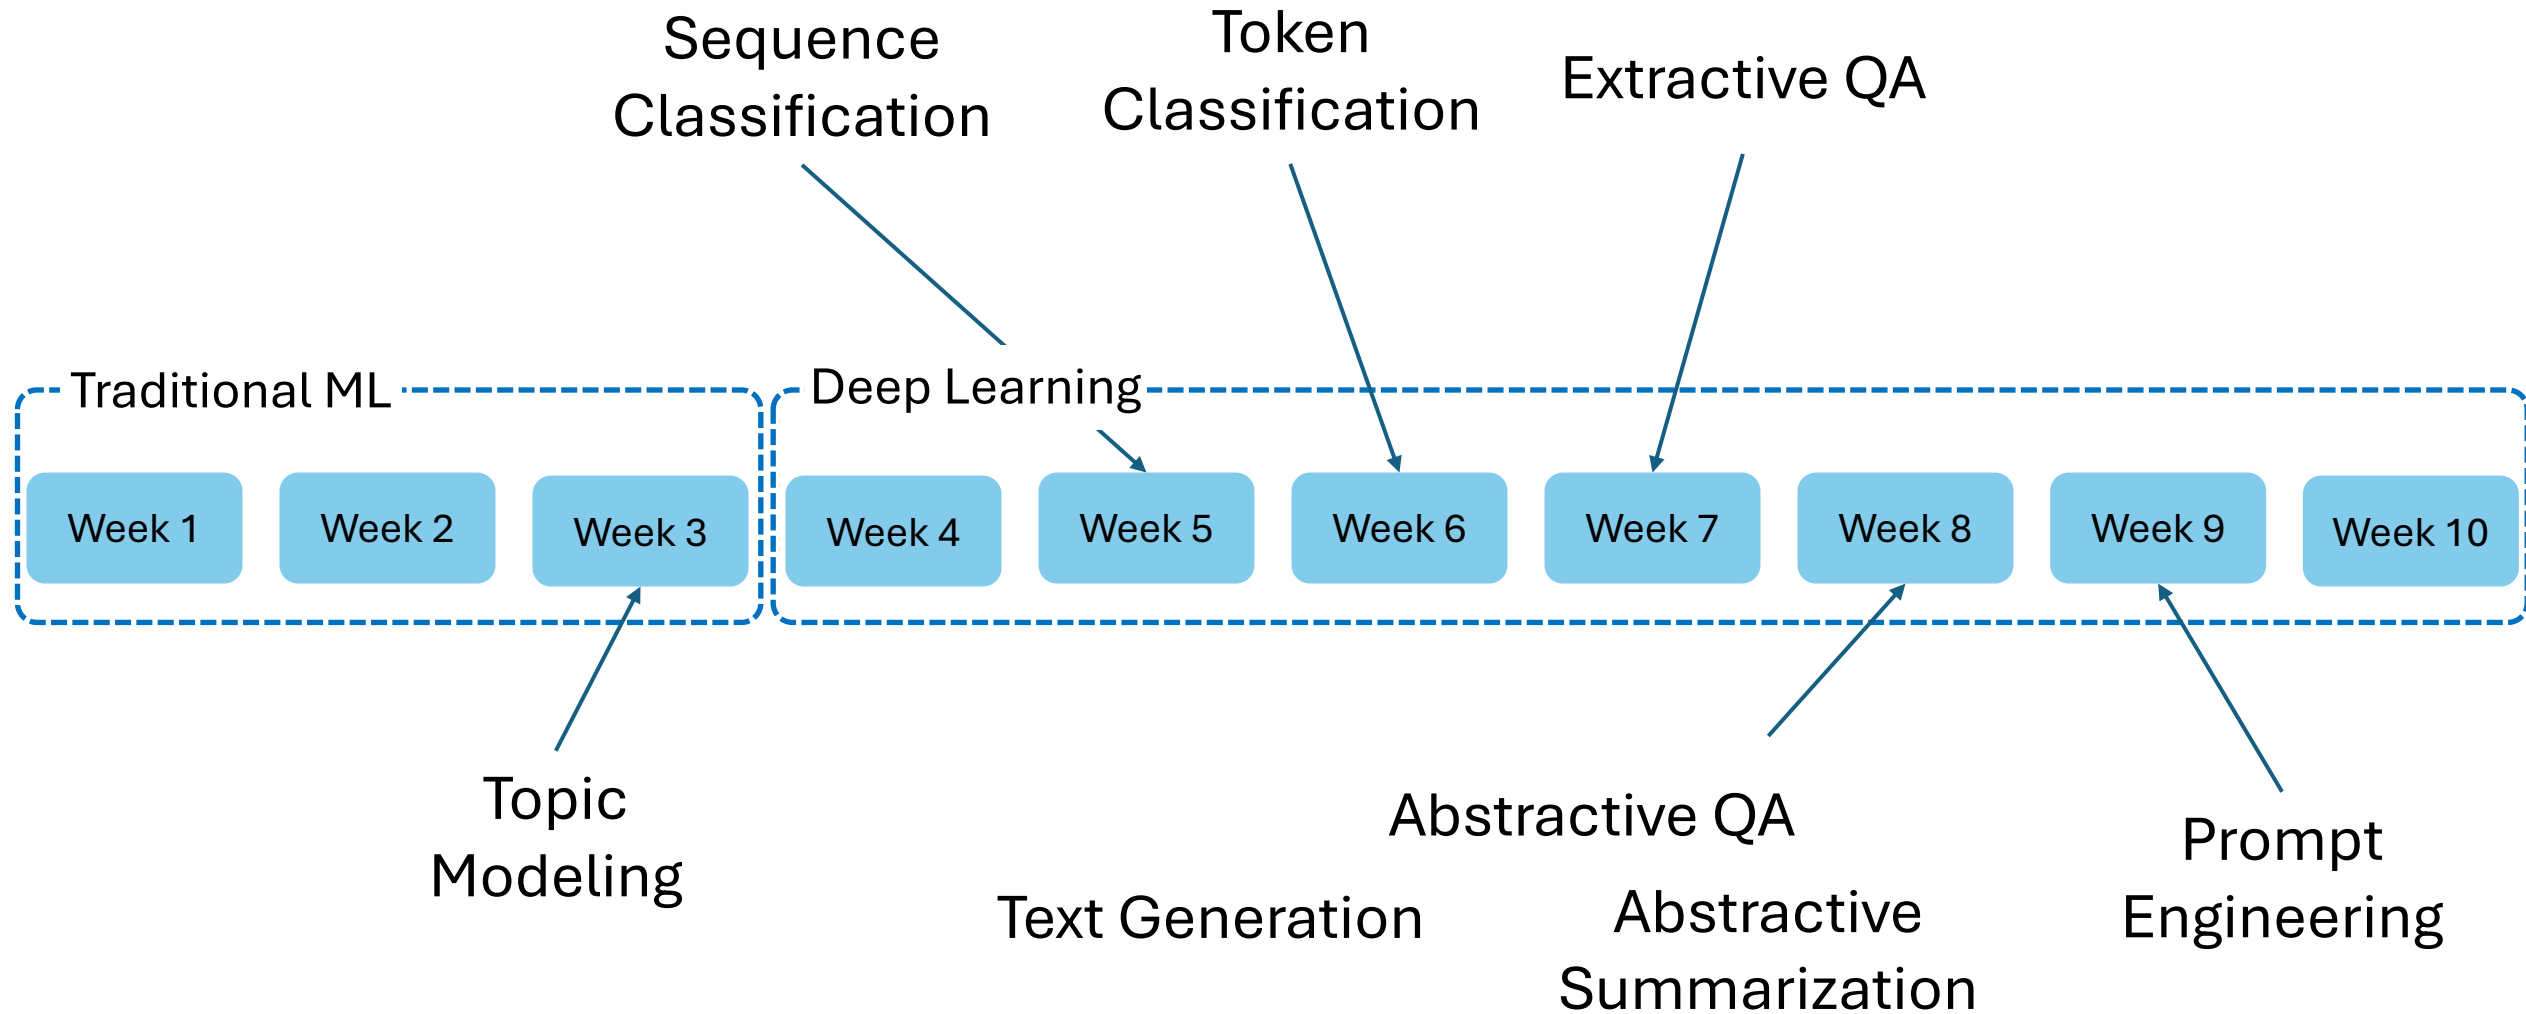


Ignas Grabauskas ([ignasg@uchicago.edu](mailto:ignasg@uchicago.edu))  
Data Scientist @ Simpson Thacher Bartlett  
Lecturer @ UChicago Data Science Institute

# Student Introductions

- Name
- Part-time? Full-time?
- Education / work background
- How far along you are in the program
- What you hope to get out of the class

# Course Introduction



# Syllabus Highlights

- If contacting instructor, TA, or grader please use email, **not** Canvas messages or comments
- Please bring laptops to class
  - **Create a python environment dedicated to the class**
  - Starting Week 2 there will short review quizzes
- Homework must be submitted individually
- Gen AI is allowed
- Two 24-hour grace periods are allowed for homework only. Must be requested before-hand (for other accommodations please see me).
- Absences must be requested

# Regular Expressions

# Regular Expressions (regex)

- A language for specifying patterns in text
- Useful when searching for specific patterns in text

Questions? Please do not hesitate to contact us at [support@ourcompany.com](mailto:support@ourcompany.com) (please allow 24 hours for our customer care team to respond to your questions). For more immediate service, you can also reach out to our customer care team at 1-555-555-5555. We are dedicated to ensuring your satisfaction and saving you \$\$\$.

```
[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}
```



# Character Matching

- Most characters will simply match themselves literally (keep case sensitivity in mind)

Questions? Please do not hesitate to contact us at support@ourcompany.com (please allow 24 hours for our customer care team to respond to your questions). For more immediate service, you can also reach out to our customer care team at 1-555-555-5555. We are dedicated to ensuring your satisfaction and saving you \$\$\$.

"Questions"

"customer care team"

"24"

# Metacharacters ^ \$ \* + ? { } [ ] \ | . ( )

- Special characters that don't match themselves (need to be "escaped" for literal match)
- They signal some out of the ordinary thing or modify behavior in some non-literal way

Questions? Please do not hesitate to contact us at support@ourcompany.com (please allow 24 hours for our customer care team to respond to your questions). For more immediate service, you can also reach out to our customer care team at 1-555-555-5555. We are dedicated to ensuring your satisfaction and saving you \$\$\$.

"\?"

"\."

"\("

"\) "

"\\$"

# Character Classes / Character Sets [ ]

- Denote a set a characters that can be matched (character-level "or")

Behavior	Regex Pattern	Matches	Notes
List characters individually	[Hh]ello	"Hello" <b>or</b> "hello"	
Provide range of characters	[A-Za-z1-4]	Any capital letter <b>or</b> any lowercase letter <b>or</b> 1 through 4	
Negate or "compliment" a set of characters	[^A-Za-z1-4]	Any character that's <b>not</b> : capital A through Z <b>or</b> lowercase a through z <b>or</b> 1 through 4	^ character negates only if at the beginning of the set
Certain metacharacters are not active	[\$?.\]^]	\$ <b>or</b> ? <b>or</b> . <b>or</b> ] <b>or</b> ^	\ ] must be escaped ^ - may need to be escaped

# Escaping and Special Sequences: \

- Used to escape metacharacters to allow for literal matches
- Denotes a special sequence for commonly used sets of characters

`\d` matches any digit character, equivalent to `[0-9]`

`\D` matches any non-digit character, equivalent to `[^0-9]`

`\s` matches any whitespace character, equivalent to `[\t\n\r\f\v]`

`\S` matches any non-whitespace character, equivalent to `[^\t\n\r\f\v]`

`\w` matches any alphanumeric character, equivalent to `[a-zA-Z0-9_]`

`\W` matches any non-alphanumeric character, equivalent to `[^a-zA-Z0-9_]`

# Dot .

- Used to match any character except newline character (though can be forced to match newline character as well)
- Useful for situations where the exact character at a certain position is unknown

Behavior	Regex Pattern	Matches	Notes
List characters individually	.at	aat bat cat dat eat !at 9at	

# Repetition / Quantifiers { } + ? \*

- Used to specify that portions of the regex pattern must be repeated a certain number of times

Behavior	Regex Pattern	Matches	Notes
Matches 0 or more times	ca*t	ct cat caaaaat	
Matches 1 or more times	ca+t	cat caaaat	
Matches 0 or 1 times	scikit-?learn	scikitlearn scikit-learn	
Matches at least <i>m</i> and at most <i>n</i> times	\d{3,5}	123 1234 12345	{m, } no upper bound {,n} no lower bound {m} exact count

# Anchors ^ \$

- Special characters that require a pattern to be to a particular area of a string

Behavior	Regex Pattern	Matches	Notes
Matches start of string	<code>^hello</code>	" <b>hello</b> , world" but not "she said hello"	
Matches end of string	<code>world\$</code>	"hello, <b>world</b> " but not "world of wonder"	

# Grouping ( )

- Allows us to look at a sequence of characters as a unit
- Useful for extracting parts of a pattern or applying quantifiers to a group of chars

Behavior	Regex Pattern	Matches	Notes
Applies quantifier to characters as a group	(ab)+	ab abab	
Allows extraction of subsequence in pattern	(\d{3})-\d{3}-\d{4}	555-555-5555 312-555-5555	Matches phone number and can extract the area code specifically, referred to as "capturing groups". For non-capturing groups, use (?:)



# Look-ahead and look-behind ( )


- Allows us to create conditions for a match based on what follows or precedes

Behavior	Regex Pattern	Matches	Notes
Match "x" only if it is followed by "y"	Springfield(?=, IL)	" <b>Springfield</b> , IL" not "Springfield, MA"	Positive Lookahead
Match "x" only if it is not followed by "y"	Springfield(?!, IL)	" <b>Springfield</b> , MA" or " <b>Springfield</b> , MO" not "Springfield, IL"	Negative Lookahead
Match "x" only if it is preceded by "y"	(?<=Mr\. )Brown	"Mr. <b>Brown</b> " not "Dr. Brown"	Positive Lookbehind
Match "x" on if it is not preceded by "y"	(?<!Mr\. )Brown	"Dr. <b>Brown</b> " not "Mr. Brown"	Negative Lookbehind

# Alternation |

- Pattern-level "or"

Behavior	Regex Pattern	Matches	Notes
Matches several patterns	cat dog bird	"cat" <b>or</b> "dog" <b>or</b> "bird"	
Matches several characters	a c	"c" <b>or</b> "b"	Better to use character sets for character-level or
Limits the reach of the alteration	(bat bird)man	Matches "batman" <b>or</b> "birdman", <b>not</b> "bat"	



# Text Normalization

# Text Normalization

- Preprocessing text to convert it into a more uniform format for NLP applications
- Typically includes
  - Tokenizing (segmenting) words
  - Normalizing word formats
  - Segmenting sentences
  - Removing stop words
- The extent of text normalization depends on the task at hand
  - Named entity recognition relies on capitalization
  - Domain-specific scenarios

# Word Tokenization

- The task of segmenting running text into words (or subwords)  
"Behold! An example of word tokenization!"  
["Behold", "!", "An", "example", "of", "word", "tokenization", "!"]
- Punctuation as separate tokens (commas useful for parsing, periods and question marks useful for sentence boundaries, etc.)
- Intra-word punctuation: *AT&T, Ph.d., m.p.h.*
- Intra-word special characters: *\$5.99, <https://www.python.org>, #LastMinuteHomework*
- Phrases: *Rock n' Roll, Salt Lake City*

# Word Tokenization

- Language dependency
- German (noun compounds not segmented)  
Lebensversicherungsgesellschaftsangestellter
- Mandarin (no spaces between words)  
我喜欢学习

# Word Tokenization

- Top-down (rule-based) tokenization
  - Begins with a comprehensive view of language (structure, grammatical conventions)
  - Rules based on syntax, grammar, or specific patterns identified in text
  - NLTK Treebank Tokenizer or spaCy
- Bottom-up tokenization, statistics of letter sequences are used to break up words into subword tokens
  - Begins with text data and uses statistical methods to identify patterns in letter sequences
  - Byte Pair Encoding (BPE), SentencePiece, and WordPiece



# Word Normalization

- The task of putting words/tokens into a standard format
  - Reducing words to their base form (lemmatization & stemming)
  - Case folding useful for generalization in information retrieval or speech recognition)
  - Case folding not useful for sentiment analysis, information extraction (US vs us)
  - Reducing word forms (USA <--> US)
- Applications are several:
  - Make solutions more generalizable (document retrieval searching for "cars" vs "car")
  - In some cases, makes identifying signal easier

# Word Normalization: Lemmatization

- Lemma: a form of a word chosen to represent a lexeme (set words related through inflection)
- Lemmatization: the task of removing word inflections, returning a word to its lemma form

Variations	Lemmatized
running, ran, runs	run
are, was, being, were	be
He visited libraries searching for books	he visit library search for book


- Requires a sophisticated analysis of a word (grammatical tense, plurality, pos)
- Useful in cases where an accurate understanding of a word's meaning is important (because lemmatization always results in a valid word)

# Word Normalization: Stemming

- Stem: the part of the word that remains unmodified during inflection (generally)
- Crudely chopping off affixes (word elements attached to stem)

Variations	Base Form
running, ran, runs	running, ran, run
are, was, being, were	are, wa, be, were
He visited libraries searching for books	he visit librari search for book

- Porter Stemmer is designed handle wide range of irregularities without being too aggressive
- Lancaster Stemmer uses a series of 120 rules to aggressively stem a word



# Minimum Edit Distance

# Minimum Edit Distance

- Much of NLP is concerned with measuring how similar two piece of text are, including individual words
- We can compare how close two strings are to each other by calculating minimum edit distance
- **Minimum edit distance:** minimum number of edit operations (insertion, deletion, substitution) needed to transform one string to another

# Alignment

I N T E \* N T I O N

| | | | | | | | |

\* E X E C U T I O N

d s s i s

**Levenshtein**

1 1 1 1 1

$$\Sigma = 5$$

**Weighted Levenshtein**

1 2 2 1 2

$$\Sigma = 8$$

# Variations

- There are different variations of minimum edit distance, for different applications
- **Damerau-Levenshtein:** allows transposing of adjacent characters. Useful for typographical errors
- **Weighted Levenshtein:** assigns different cost to different operations. Useful for keyboard layout optimization or OCR
- **Jaro-Winkler:** gives favorable scores to strings that match from the beginning. Useful for short names and record linkage
- And so on...



